Backstory Consistency Verification Report

Task: Binary classification of backstory consistency with novel
Method: Claim -> Evidence -> Verdict using RAG-based retrieval
Sample: Kingdom of Eldoria novel and Prince Adrian backstory

Results Summary:

Total Claims Evaluated: 4
Verdicts: 3 SUPPORT, 0 CONTRADICT, 1 UNCLEAR
Final Classification: 0 (Contradictory - one claim unclear)

Pipeline Components:

1. Text Chunking: 500-char chunks with 100-char overlap
2. Evidence Retrieval: Keyword-based similarity (top-5 chunks)
3. Claim Evaluation: Heuristic-based (LLM-ready placeholder)
4. Aggregation: Conservative logic (any non-support = 0)

Key Metrics:

- Handles 100k+ word novels without full context to LLM
- ~5 chunks per claim (~2.5k tokens max)
- No model training. Conservative verdicts.

Output:

- Binary label: 0 or 1
- JSON report with per-claim verdict and evidence count
- CSV results table